

# Unsupervised Learning Notes

## Contents

<b>Introduction to unsupervised learning</b>	<b>2</b>
Idea of unsupervised learning . . . . .	2
Readings . . . . .	3
lab . . . . .	3
<b>Distance and similarity</b>	<b>3</b>
Feature scaling and normalization . . . . .	3
Distance metrics . . . . .	4
Curse of dimensionality . . . . .	6
Readings . . . . .	6
lab . . . . .	6
<b>Clustering methods</b>	<b>6</b>
Intro to clustering . . . . .	6
Partitioning methods: K-means and k-medoids clustering . . . . .	7
Hierarchical Methods . . . . .	10
Density-based spatial clustering of applications with noise (DBSCAN) . . . . .	12
Comparison of k-means, hierarchical, and DBSCAN methods . . . . .	14
Assessing cluster quality . . . . .	14
Readings . . . . .	15
lab . . . . .	16

# Introduction to unsupervised learning

## Idea of unsupervised learning

1. Slideshow: What is AI, ML, data science, statistics, etc? How does it all fit? Cool diagrams.
2. Supervised learning vs unsupervised learning:
  - (a) Supervised learning:
    - Given labeled data, predict a target. Example is predicting a student's GPA from student data.
    - More formally, labeled training data is  $(X, y)$  for  $n$  samples with relationship

$$y = f(X) + \epsilon.$$

Goal is to approximate  $f$  via  $\hat{f}$ . Call supervised because we know answers from our training data.

- Basic example of linear regression with  $p = 1$  for student data. Assume  $f$  is linear. Graph with data and best fit line  $\hat{f}$ . Write general  $p$  variable case. Can be used for student intervention (early warning, effect of living on campus, attending class).
- Main goals are prediction and inference by understanding  $f$  (feature importance, model fit and reliability).
- Linear model is an assumption. Other models include decision trees, neural networks, support vector machine, and more.
- Well understood area, clear ways to assess quality of results. ISLR text key reference.

- (b) Unsupervised learning:

- Unsupervised learning is a class of machine learning methods used to discover structure in data

$$X = \{x_1, x_2, \dots, x_n\}, \quad x_i \in \mathbb{R}^p \quad (\text{note vector notation})$$

without labeled outcomes. Instead of predicting a known target, the goal is to explore, summarize, and reveal patterns that are intrinsic to the data.

- Given unlabeled data  $X$ , find structure in the data. No prediction, no supervision. Learning by observation, not by example.
- What types of students are there given hours studied per week and GPA? What variable combinations belong together (academic, engagement)? Unusual students? Goal is to better understand data.
- Can be a stand alone analysis or can be used to compliment supervised learning.

3. Core tasks of unsupervised learning:

- (a) Clustering: Group similar observations in the same cluster. K-means, hierarchical, dbscan
- (b) Dimension reduction: Reduce noise and multicollinearity, data viz, large to smaller data, data understanding. PCA, t-SNE, UMAP, SVD. Google tensorflow embedding projector.
- (c) Anomaly detection: Learn distribution of data to quantify outlier probabilities.

4. Concrete examples of unsupervised learning:

- (a) Customer segmentation: Clustering
  - Walmart data on spend average, frequency, mode, product mix, app use.

- No label such as budget shopper or family provider. Want to discover segments rather than predetermine behavior.
- Each data point is a customer in high dimensions.
- Similarity means close distance. Scaling and choice of distance matters.
- Possible clusters: Weekly family stockup, single essentials, deal hunters. These are business driven interpretations.
- Actions: Store layout optimization, personal coupons, inventory planning, regional differences. Not aiming for individual predictions (as with supervised learning).
- Google: Walmart customer segmentation

(b) Spotify music genre: Dimension reduction

- Google: Spotify api dataset
- Many automatic features, how to tell what genre?
- Vectors are high dimension, but human perception is low dimension.
- Can we compress data into low-dimensions?
- Distance reflect song similarity.
- Are there distinct groups of genres or continuous flow?
- Goal is to make data more intelligible.

(c) Credit card fraud: Anomaly detection

- Each data point is a transaction (amount, time, source, location, recent freq, device).
- Millions of these per day, tiny amount are fraud.
- False positive is a problem.
- Does this deviate from expected? Normal behavior is dense regions, but some deviance can naturally occur.
- Distance metric determines deviation from normal.
- Per customer (card?) normalization.
- Action may be to identify fraud patterns to catch more.

5. Much more challenge than supervised learning. No simple goal. Results are subjective. Exploratory, descriptive, and hypothesis-generating rather than predictive.

## Readings

1. ISLR 2.1.4, Ch12 thru 12.1
2. HOUL Ch1

## Lab

1. EDA and data cleaning, DMCT Ch2 and Ch3

## Distance and similarity

### Feature scaling and normalization

1. Features on a bigger scale dominate distance calculations. Spend vs weekly visit count for customer segmentation. Illustrate dollars vs thousands for spend on Euclidean distance.

$$d_2(x, y) = \|x - y\|_2 = \sqrt{\sum (x_j - y_j)^2}$$

Many ways to scale to prevent this issue.

## 2. Standardization ( $z$ -score scaling):

- (a) Definition: Shift by mean and divide by standard deviation.

$$x'_i = \frac{x_i - \bar{x}}{s_x}$$

where  $s_x$  is the standard deviation of variable  $x$ . Remind of standard deviation calculation. Recall sample vs population notation. Draw distribution picture, two different shapes.

$$s_x = \sqrt{\frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2}, \quad \bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$$

- (b) Results in mean 0, variance 1. Now on a standard normal distribution  $N(0, 1)$ . Nice properties of a normal distribution for 68-95-99 percent confidence intervals. Preserves order and relative distribution.
- (c) Equal feature contribution, preserves relative difference.
- (d) Extensions include mean absolute deviation.

## 3. Min-max scaling:

- (a) Definition: Shift by min and divide by range.

$$x'_i = \frac{x_i - \min(x)}{\max(x) - \min(x)}$$

- (b) Guess what it does. Refer to mall data scaling age.
  - (c) Result is in interval  $[0, 1]$ . Preserves order but compresses extremes.
  - (d) Good for bounded features.
- 4. Most of the time scale, can always reverse scaling if you keep track of original stats. When scaling is NOT a good idea:
    - (a) Feature units are meaningful, important, and want to keep for analysis.
    - (b) Binary indicators (0/1), changes to discrete values but different values.
    - (c) Counts with semantic meaning, such as a 1-10 satisfaction rating.
    - (d) Ratios which are already normalized.

## Distance metrics and similarity

- 1. Distance is a choice we make which encodes different notions of similar. Many options.
- 2. Key properties are required: Mathematically called a metric in a metric space.
  - (a) Positivity:  $d(x, y) \geq 0$  for all  $x, y$  and  $d(x, y) = 0$  only if  $y = x$ .
  - (b) Symmetry:  $d(x, y) = d(y, x)$ .
  - (c) Triangle inequality:  $d(x, z) \leq d(x, y) + d(y, z)$  for any  $x, y, z$ . Illustrate with pictures,  $x$  and  $z$  on the horizontal axis.
- 3. Common distance metrics:

(a) Euclidean distance ( $\ell_2$  norm): For points  $x, y \in \mathbb{R}^d$ ,

$$d_2(x, y) = \|x - y\|_2 = \sqrt{\sum (x_i - y_i)^2}$$

- Geometry: Straight-line distance, rotation invariant, penalizes large feature deviations heavily.
- Assumes: Features are commensurate, spherical neighborhoods make sense
- Example: Customers close if very same spend pattern and volume. (same shop and spend)

(b) Manhattan distance ( $\ell_1$  norm):

$$d_1(x, y) = \|x - y\|_1 = \sum |x_i - y_i|$$

- Geometry: City block distance, diamond shaped contours, less sensitive to large feature deviations.
- Assumes: More robust to outliers.
- Example: Customers close if same spend pattern and volume, some diff tolerated. (similar with occasional deviation allowed)

(c) Supremum ( $\ell_\infty$  norm):

$$d_\infty(x, y) = \max\{|x_i - y_i|\}$$

- Geometry: Largest difference only. Shapes are squares.
- Example: Customers close if same spend pattern and volume, one big difference is a problem.

(d) These are all generalized by Minkowski distance:

$$d(x, y) = \|x - y\|_r = \sqrt{\sum (x_i - y_i)^r}$$

Overall, Euclidean is the favorite, but there are good reasons to use others. Show animation in Desmos of different unit circles  $x^r + y^r = 1$ .  $r = \infty$  case is the limit which calc 1 shows supremum norm is found. Show in 2D.

(a) Similarities: Note not a distance, but sometimes makes more sense.

(b) Cosine similarity:

$$\text{sim}_{\cos}(x, y) = \frac{x \cdot y}{\|x\| \|y\|} = \cos(\theta)$$

- Geometry: measures cosine of angle between rather than distance, lives in  $[-1, 1]$ , same direction is close to 1, orth and opposite direction, ignores scale. Turns our ordering agrees with Euclidean distance.
- Assumes: Pattern rather than intensity.
- Example: Customers close if same spend pattern but different spend volumes. (shop same regardless of spend)

(c) Correlation: Sample Pearson correlation.

$$\text{sim}_{\text{cor}}(x, y) = r = \text{cor}(x, y) = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}}$$

- Geometry: Measures linear relationship,  $r \approx 1$  strong positive linear relationship,  $r \approx -1$  strong negative, in between. Book picture vs Wikipedia picture.
- Same as cosine similarity if data is normalized.

## Curse of dimensionality

1. In high dimensions, all points become almost equally far apart. Nearest points are almost same distance as farthest.
2. As dimension  $d$  increases,

$$\frac{\max d(x, y) - \min d(x, y)}{\min d(x, y)} \rightarrow 0$$

## Readings

1. IDM Ch 2 Intro, 2.3.7, 2.4.1-2.4.5, 2.4.9-2.4.10 recommend all of Ch 2
2. DMCT Ch 2 Intro, 2.5.1, 2.3.1, 2.3.4, 2.3.7, 2.3.9 recommend all of Ch 2 except 2.6

## Lab

## Clustering methods

Main reference: DMCT chapter 9

### Intro to clustering

1. What is cluster analysis?
  - (a) Clustering groups data into clusters so that objects within a cluster are more similar to each other than to objects in other clusters, according to a chosen notion of similarity.
  - (b) Clustering is an *ill-posed* problem: there is no single correct solution. Results depend on modeling assumptions, similarity definitions, and analysis goals. There is no universal notion of “true” clusters.
  - (c) Similarity (or dissimilarity) is typically defined through a distance or proximity measure.
  - (d) Different clustering methods reflect different assumptions about data structure (e.g., partitioning, hierarchical, density-based, grid-based). Some methods that optimize an explicit objective (k-means) while others identify structure without a global objective (hierarchical, DBSCAN).
  - (e) Clustering quality can be assessed in multiple ways, including internal criteria, external validation, stability, and interpretability.
  - (f) Ongoing research focuses on scalability, high-dimensional settings where distance metrics break, complex cluster shapes, and diverse data types (e.g., text, images).
2. Desiderata for clustering methods:
  - (a) Ability to handle different data types (numeric, categorical, mixed)
  - (b) Ability to detect non-spherical or non-convex clusters
  - (c) Robustness to noise and outliers
  - (d) Scalability to large datasets
  - (e) Ability to incorporate constraints or side information
  - (f) Results that are interpretable and actionable
3. Types of clustering methods

- (a) Partitioning methods (customer segmentation, geog segm)
  - Assume clusters are compact, well-separated, and cover all data points.
  - Partition  $n$  objects into  $k$  non-overlapping clusters.
  - Typically distance-based and solved via iterative optimization.
  - Sensitive to initialization, distance choice, and cluster shape assumptions.
- (b) Hierarchical methods (fish in Mississippi, species, genus, family, order, class, phylum, kingdom)
  - Assume nested cluster structure is meaningful.
  - Produce a hierarchy of clusters represented as a tree.
  - Agglomerative (bottom-up) or divisive (top-down) approaches.
  - Once a merge or split occurs, it cannot be undone.
- (c) Density-based methods (social data with natural crowding, maybe no fixed  $k$  or hierarchy)
  - Assume clusters correspond to regions of high data density separated by low-density regions.
  - Clusters are grown based on neighborhood density criteria.
  - Naturally identify outliers and allow arbitrary cluster shapes.
- (d) Key challenges and limitations:
  - Evaluation is inherently difficult due to lack of ground truth; metrics often encode the same assumptions as the algorithm.
  - Clustering is exploratory rather than confirmatory.
  - Domain knowledge plays a central role (feature selection, scaling, similarity choice, interpretation, clustering alg choice).
  - Clustering is not classification, causal inference, or discovery of objective real-world categories; results should not be over-interpreted.

## Partitioning methods: K-means and k-medoids clustering

### 1. Big picture:

- (a) Given data  $D = \{x_1, \dots, x_n\} \subset \mathbb{R}^p$ , partition into  $k$  disjoint clusters  $C_1, \dots, C_k$ ,  $C_i \cap C_j = \emptyset$ .
- (b) Each point belongs to exactly one cluster, and each cluster is summarized by a single representative (center).
- (c) Clustering defined via optimization of an objective function called within-cluster loss.

### 2. Key assumptions:

- (a) A meaningful distance  $d(x, y)$  exists (eg Euclidean distance).
- (b) Clusters are compact and well-separated in the chosen metric.
- (c) Cluster centers summarize cluster geometry.
- (d) All data belong to some cluster (no noise model).
- (e)  $k$  is fixed and meaningful.

### 3. k-means: A centroid-based technique

- (a) Cluster center: centroid (mean)  $c_i \in \mathbb{R}^p$ .

(b) Objective function to minimize:

$$\min_{C_1, \dots, C_k} \sum_{i=1}^k \sum_{x_j \in C_i} d(x_j, c_i)^2, \quad c_i = \frac{1}{|C_i|} \sum_{x_j \in C_i} x_j.$$

where the distance metric is Euclidean distance (almost always)

$$d(x, y) = \|x - y\| = \sqrt{(x_1 - y_1)^2 + \dots + (x_p - y_p)^2}.$$

The centroid represents a hypothetical average cluster member (customer).

(c) The within cluster variance is the sum of square errors:

$$WCSS = \sum_{i=1}^k \sum_{x_j \in C_i} d(x_j, c_i)^2$$

(d) Algorithm (Lloyd's algorithm):

- This optimization problem is computationally expensive (NP-hard), many ways to select  $k$  clusters of  $n$  data points. A basic iterative algorithm is used instead.
- Initialize  $c_1, \dots, c_k$  as  $k$  random objects from  $D$ .
- Cluster assignment step: Belong to nearest centroid.

$$x_j \mapsto \arg \min_i \|x_j - c_i\|^2.$$

- Centroid update step: Average of cluster members.

$$c_i \leftarrow \frac{1}{|C_i|} \sum_{x_j \in C_i} x_j.$$

- Iterate until assignments stabilize.

- (e) Each step decreases the objective, converges to a local minimum, may not be global minimum.
- (f) Implies Voronoi partition of the feature space, boundaries between clusters where two centroids are equidistant, only depends on centroid not data distribution.
- (g) Implicit assumptions: spherical clusters, equal variance, Euclidean geometry.
- (h) Variations of  $k$ -means involved different distance metrics, smart centroid initialization, and centroid calculation strategies,  $k$ -modes for nominal data, groupings of data called microclusters.

#### 4. $k$ -medoids: A representative object-based technique

- (a) Motivation:  $k$ -means is sensitive to outliers when centroids (means) are calculated.
- (b) Cluster center: medoid  $o_i \in x_1, \dots, x_n$ . Centroid is now a data point. Also called a representative object.
- (c) Objective function to minimize: Note that medoids  $o$  must be data points, so computing new centers is not obvious.

$$\sum_{i=1}^k \sum_{x_j \in C_i} d(x_j, o_i)$$

Distance  $d(\cdot, \cdot)$  need not be Euclidean. Note the lack of squared distance.

(d) More robust to outliers (no averaging).

(e) Example algorithm: Partitioning around medoids PAM

- Random initial medoids. Assign cluster membership. Brute force check all updated medoids per assigned cluster to see which minimizes total distance error. Iterate.

$$WCTD = \sum_{i=1}^k \sum_{x_j \in C_i} d(x_j, o_i)$$

- Modification for large data, clustering large applications (CLARA) considers random samples of the dataset.

(f) Slower than  $k$ -means due to discrete optimization, especially for large  $n$  and large  $k$ .

#### 5. $k$ -means vs $k$ -medoids (mathematical contrast)

- Continuous optimization (means) vs discrete optimization (medoids).
- Squared Euclidean loss vs general metric loss.
- Sensitive vs robust to outliers.
- Fast gradient-like updates vs combinatorial search.

#### 6. Practical issues:

- Objective function is typically non-convex  $\Rightarrow$  multiple local minima. Repeat iterations can give different cluster solutions.
- Initialization matters (e.g., random vs  $k$ -means++ which chooses initial clusters far apart).
- Scaling changes the geometry of  $|\cdot|$ .
- Choice of  $k$  is a modeling decision, not a statistical estimate. If you don't know  $k$ , try many and compare results.

#### 7. When partitioning methods work well

- Clusters roughly convex and isotropic (not stretched in a certain direction).
- Moderate dimension with meaningful distances, otherwise dimension reduction needed.
- Clear notion of "center."
- Need fast baseline clustering.

#### 8. When they fail

- Non-convex or nested clusters.
- Unequal cluster variances or densities.
- Strong outliers (especially  $k$ -means).
- High-dimensional distance concentration.

#### 9. How to determine the right number of clusters? Many ways! Basic discussion here, more later.

- Need  $k$  in advance, but the right number of clusters is often ambiguous.
- Good to balance compressibility (simplifying data) with accuracy (cluster meaning and usability).
- An easy approach considers plotting WCSS (within cluster sum of squares) for  $k = 2, 3, 4, \dots, k$  and use the elbow method. Sharp bend says smaller improvements. Expect to decrease to 0. Subjective but something to go by. Better(?) ways!

#### 10. Fun simulator: <https://clustering-visualizer.web.app/kmeans>

## Hierarchical Methods

### 1. Big picture:

- (a) Hierarchical clustering builds a nested sequence of partitions.
- (b) Output is a tree (dendrogram), not a single clustering.
- (c) Clusters exist at multiple resolutions (choices of  $k$ ).
- (d) No single "best" number of clusters is assumed a priori.
- (e) 2 main types: Agglomerative (bottom-up) and divisive (top-down)

### 2. Key assumptions:

- (a) Nested structure in the data is meaningful.
- (b) Pairwise dissimilarities capture relevant structure.
- (c) Early decisions (merges or splits) are trustworthy.
- (d) No noise model: all points participate in the hierarchy.

### 3. Agglomerative hierarchical clustering:

#### (a) Algorithm: Bottom-up procedure.

- Start with  $n$  singleton clusters  $\{x_1\}, \{x_2\}, \dots, \{x_n\}$ . Compute a pairwise distance matrix.
- Iterative merge the two closest clusters. The dissimilarity between these two clusters indicates the height in the dendrogram where the fusion is placed.
- Compute updated pairwise inter-cluster dissimilarities among the  $n - 1$  remaining clusters. Repeat merging.
- Merge low in dendrogram means joined clusters were similar. High merge means clusters were dissimilar (separation).

#### (b) Cluster-cluster distance (linkage)

- Requires a linkage function  $D(C_a, C_b)$  for clusters  $C_a, C_b$ .
- Common choices:

$$\text{Single / Minimum / Nearest neighbor: } D(C_a, C_b) = \min_{x \in C_a, y \in C_b} d(x, y)$$

Allows long, thin, winding, non-convex clusters. Connectivity matters, global compactness does not.

$$\text{Complete / Maximum / Farthest neighbor: } D(C_a, C_b) = \max_{x \in C_a, y \in C_b} d(x, y)$$

Produces compact, spherical clusters and penalizes irregular shapes. Sensitive to outliers.

$$\text{Average: } D(C_a, C_b) = \frac{1}{|C_a||C_b|} \sum_{x \in C_a} \sum_{y \in C_b} d(x, y)$$

$$\text{Mean / Centroid: } D(C_a, C_b) = |m_a - m_b|$$

Average and mean are compromise between single and complete linkage.

- (c) Linkage choice encodes shape assumptions. Distance metric also important per above discussion (what similar should mean in application).

### 4. Ward's method (variance-based linkage)

- Merge clusters that minimally increase total within-cluster variance.
- Objective interpretation:

$$\begin{aligned} W(C_i, C_j) &= WCSS(C_i \cup C_j) - WCSS(C_i) - WCSS(C_j) \\ &= \sum_{\vec{x} \in C_i \cup C_j} d(\vec{x}, \vec{c}_{ij})^2 - \sum_{\vec{x} \in C_i} d(\vec{x}, \vec{c}_i)^2 - \sum_{\vec{x} \in C_j} d(\vec{x}, \vec{c}_j)^2 = \frac{|C_i||C_j|}{|C_i| + |C_j|} \|\vec{c}_i - \vec{c}_j\|^2 \end{aligned}$$

where  $c$  denotes the cluster centroid as in k-means. Note only for Euclidean distance for last step.

- Closely related to k-means objective, though indirectly. Best of both worlds in a way.
- Favors compact, spherical clusters.

## 5. Dendrogram: Illustrate example with basic distance measures.

- Tree structure encoding merge order and merge distances.
- Vertical height = dissimilarity at which merge occurs. Note near in the horizontal direction does not mean points/clusters are near. Good discussions in ISLR,
- Cutting the tree at height  $h$  induces a partition.
- Different cuts correspond to different  $k$ .

## 6. How to decide cut?

- Unlike k-mean, HCA does not optimize a cost function such as WCSS. Other approaches are needed.
- Dendrogram viz is key.
- Elbow method of linkage distance vs merge step (A sharp increase in linkage distance = merging dissimilar clusters, cut just before the big jump)
- WCSS works for Ward's linkage because it is inherently close to  $k$ -means.
- Cophenetic Correlation Coefficient (CCC)
  - The cophenetic correlation coefficient (CCC) is specific to hierarchical clustering and is often overlooked.
  - Idea: How well does the dendrogram preserve the original pairwise distances?

$$CCC = \text{corr}\left(d_{ij}, \hat{d}_{ij}\right)$$

where  $d_{ij}$  is the original distance between points  $i$  and  $j$ , and  $\hat{d}_{ij}$  is the cophenetic distance, defined as the height at which points  $i$  and  $j$  merge in the dendrogram.

- Interpretation: Values close to 1 indicate that the hierarchical clustering preserves pairwise distances well. Low values indicate that the dendrogram substantially distorts the geometry of the data.
- Uses: Comparing linkage methods (single, complete, average, Ward), and choosing an appropriate distance metric.

## 7. Divisive hierarchical clustering

- Top-down approach.
- Start with all points in one cluster.
- Recursively split clusters.

- Less common due to computational cost.
  - Conceptually closer to repeated partitioning.
8. When hierarchical clustering works well
- Data have meaningful nested or multi-scale structure.
  - Moderate sample size.
  - Interest in relationships between clusters, not just assignments.

9. When it fails

- Large datasets (computational and memory cost).
- Strong noise or chaining effects (single linkage).
- Early incorrect merges propagate upward.
- Noisy distance measurements.

### Density-based clustering: DBSCAN

1. Big picture:
  - (a) DBSCAN = Density-based spatial clustering of applications with noise
  - (b) Clusters are defined as regions of high point density separated by regions of low density.
  - (c) Does not impose global geometry (no centroids, no partition).
  - (d) Explicitly allows noise and outliers.
  - (e) Number of clusters is determined by the data, not fixed in advance.
2. Key assumptions:
  - (a) A meaningful distance metric exists.
  - (b) Clusters correspond to dense regions in the metric space.
  - (c) Density is approximately homogeneous within the clusters.
  - (d) Low-density regions separate clusters.
3. Parameters:
  - (a)  $\varepsilon > 0$  (radius parameter).
  - (b)  $\text{minPts} \in \mathbb{N}$  (minimum number of neighbors)
4. Neighborhood definition:
  - (a)  $\varepsilon$ -neighborhood of a point  $x$ :
$$N_\varepsilon(x) = \{y : d(x, y) \leq \varepsilon\}$$
  - (b)  $|N_\varepsilon(x)|$  gives the density of point  $x$ , including the point itself in this count.
5. Point types:
  - (a) Core point: At least  $\text{minPts}$  within a neighborhood of that point. Point  $A$  in diagram if  $\text{MinPts} \geq 7$ .
$$|N_\varepsilon(x)| \geq \text{minPts}$$

- (b) Border point: Not core, but falls in the neighborhood of a point (within the border), or possibly many points. Point  $B$  in diagram.

$$|N_\varepsilon(x)| < \text{minPts}, \quad \text{but } x \in N_\varepsilon(y) \text{ for some core point } y$$

- (c) Noise point:  $x$  is neither core nor border. Point  $C$  in diagram.

6. Density reachability:

- (a) Directly density-reachable:

$$y \in N_\varepsilon(x), \quad x \text{ is a core point}$$

- (b) Density-reachable: chain of directly density-reachable points.

- (c) Density-connected: two points reachable from a common core point.

7. Cluster definition:

- (a) A cluster is a maximal set of density-connected points.

- (b) Noise points are not assigned to any cluster.

8. Algorithm (conceptual):

- (a) Identify all core points by checking neighborhood density of all possible points.

- (b) Grow clusters by connecting density-reachable points into the same cluster.

- (c) Label remaining points as noise or border.

9. Geometric consequences:

- (a) Can recover non-convex and arbitrarily shaped clusters.

- (b) No forced assignment of all points.

- (c) Cluster boundaries follow low-density regions.

- (d) No global partition space.

10. Comparison to partitioning methods

- (a) No centroids or objective function.

- (b) No Voronoi geometry.

- (c)  $k$  not specified.

- (d) Explicit noise handling.

11. Sensitivity and limitations

- (a) Choice of  $\varepsilon$  and minPts is critical.

- (b) Struggles with varying cluster densities.

- (c) Distance concentration in high dimensions degrades performance.

- (d) Sensitive to distance scaling.

12. When DBSCAN works well

- (a) Clusters separated by low-density regions.

- (b) Non-spherical, irregular shapes.

- (c) Presence of noise or outliers.
  - (d) Low-to-moderate dimensional data.
13. When it fails
- (a) Clusters with significantly different densities.
  - (b) High-dimensional data.
  - (c) Data without clear density gaps.
  - (d) Poorly chosen distance metric.

### Comparison of k-means, hierarchical, and DBSCAN methods

1. Key idea: Clustering methods are not interchangeable algorithms. They encode fundamentally different notions of what a cluster is.
2. The question each method answers:
  - (a) Partitioning (k-means and k-modes): Given  $k$ , how should I divide all points to minimize within-cluster dissimilarity?
  - (b) Hierarchical clustering: How are points related across multiple scales of similarity?
  - (c) DBSCAN: Which points belong to the same dense region, and which points are noise?
3. Mathematics of the machines:
  - (a) Partitioning:
    - Force the data into  $k$  compact clusters by minimizing within-cluster loss
    - Imposes Voronoi partition where cluster boundaries are hyperplanes
    - Parameters are  $k$ , distance metric, scaling, initialization
    - Fails for non-convex shapes, unequal variance, noise
  - (b) Hierarchical:
    - Reveal nested similarity structure through greedy merges or splits
    - No global geometric partition and shape depends on linkage choice
    - Parameters are distance metric, linkage, cut height
    - Fails for noise, chaining, large  $n$
  - (c) DBSCAN:
    - Identify dense regions separated by low-density gaps and label the rest as noise
    - No partition of space and geometry adapts to data distribution
    - Parameters are  $\varepsilon$ , minPts, distance metric
    - Fails for varying densities, high dimension

### Assessing cluster quality

1. Fundamental difficulty:
  - (a) Clustering is unsupervised: typically no ground truth labels, may not even be clusters in the first place.
  - (b) "Good clustering" is not uniquely defined. Many metrics.
  - (c) Evaluation criteria often encodes the same assumptions as the algorithm, can be deceptive.

- (d) Different metrics may rank the same clustering very differently.  
(e) Validation is about usefulness and stability, not correctness.
2. Three perspectives on validation:
- (a) Internal validation: use only the data and clustering structure
  - (b) External validation: compare to known labels (when available)
  - (c) Relative validation: compare multiple clusterings to each other
3. Internal validation (geometry-based):
- (a) Measures compactness (within-cluster similarity) and separation (between-cluster dissimilarity).
  - (b) Common quantities:
- $$W = \sum_{k=1}^K \sum_{i \in C_k} \|x_i - \mu_k\|^2 \quad (\text{within-cluster dispersion})$$
- $$B = \sum_{k=1}^K n_k \|\mu_k - \mu\|^2 \quad (\text{between-cluster dispersion})$$
- These favor spherical, equal-variance clusters.
- (c) Silhouette analysis:
- For point  $i$ :
- $$a(i) = \text{average distance to points in same cluster}$$
- $$b(i) = \min_{k \neq c(i)} \text{average distance to cluster } k$$
- $$s(i) = \frac{b(i) - a(i)}{\max\{a(i), b(i)\}} \in [-1, 1]$$
- $s \approx 1$ : well-clustered
  - $s \approx 0$ : ambiguous
  - $s < 0$ : likely misclustered
  - Assumes distance-based, compact clusters.
- (d) Silhouette is an internal validation metric like  $W$  and  $B$ , but it evaluates clustering locally at the point level rather than globally at the centroid level.
- 4.

## Readings

1. DMCT Ch10, cluster analysis basic concepts and methods
2. DMCT 10.1, intro to cluster analysis
3. DMCT 10.2, partition methods (k-means and k-medoids)
4. ISL 12.4.1, k-means clustering
5. DMCT 10.3, hierarchical methods
6. ISL 12.4.2, hierarchical clustering

7. DMCT 10.4.1, DBSCAN
8. ISL 12.4.3, practical issues in clustering
- 9.

**Lab**